

Shared Memory Programming




(I)Shared Memory Programming Model



Shared memory programming with threads.

▣ Global memory




▣ Parallelization by threads

run time
→


Master thread
runs in serial
mode

Shared memory parallel programming with threads

▣ Global memory




▣ Parallelization by threads




Shared memory parallel programming with threads

▣ Global memory




▣ Parallelization by threads



Shared memory parallel programming with threads


Multiple workers at the same shared white board or workspace.



(II) Heat Equation by Finite Differences



Simple example problem: temperature profile in a rod




- ▣ Rod heated uniformly by burners, ends held at constant temperature.
- ▣ Idealized equations

$$\frac{\partial u(x,t)}{\partial t} - \frac{\partial^2 u(x,t)}{\partial x^2} = 1; \quad 0 < x < 1; \quad 0 < t < 2$$

$$u(0,t) = u(1,t) = 0; \quad u(x,0) = 0$$

Simple example problem: temperature profile in a rod [Steady State]




- ▣ Rod heated uniformly by burners, ends held at constant temperature.
- ▣ Idealized equations

$$\cancel{\frac{\partial u(x,t)}{\partial t}} - \frac{\partial^2 u(x,t)}{\partial x^2} = 1; \quad 0 < x < 1; \quad 0 < t < 2$$

Note: A red arrow points from the red '0' above the equation to the red '0' in the boundary condition below.

$$u(0,t) = u(1,t) = 0; \quad u(x,0) = 0$$

Simple example problem: temperature profile in a rod [Steady State]




- ▣ Rod heated uniformly by burners, ends held at constant temperature.
- ▣ Idealized equations

$$-\frac{\partial^2 U(x)}{\partial x^2} = 1; \quad 0 < x < 1;$$


$$U(0) = U(1) = 0;$$

Simple example problem: temperature profile in a rod [Steady State]




- ▣ Rather than looking for a solution as a continuous function, like $\sin(x)$, compute solution U at a discrete set of spatial points.

Simple example problem: temperature profile in a rod [Steady State]




- ▣ Rather than looking for a solution as a continuous function, like $\sin(x)$, compute solution U at a discrete set of spatial points.

Discretizing: going from the continuum (PDE) to the discrete (algebraic)



- ▣ Discretizing the problem also involves converting the derivatives into algebraic difference equations

Discretizing: going from the continuum (PDE) to the discrete (algebraic)



□ Taylor's theorem

$$U(x+h) = U(x) + U'(x)h + \frac{1}{2}U''(x)h^2 + \frac{1}{6}U'''(x)h^3 + O(h^4)$$

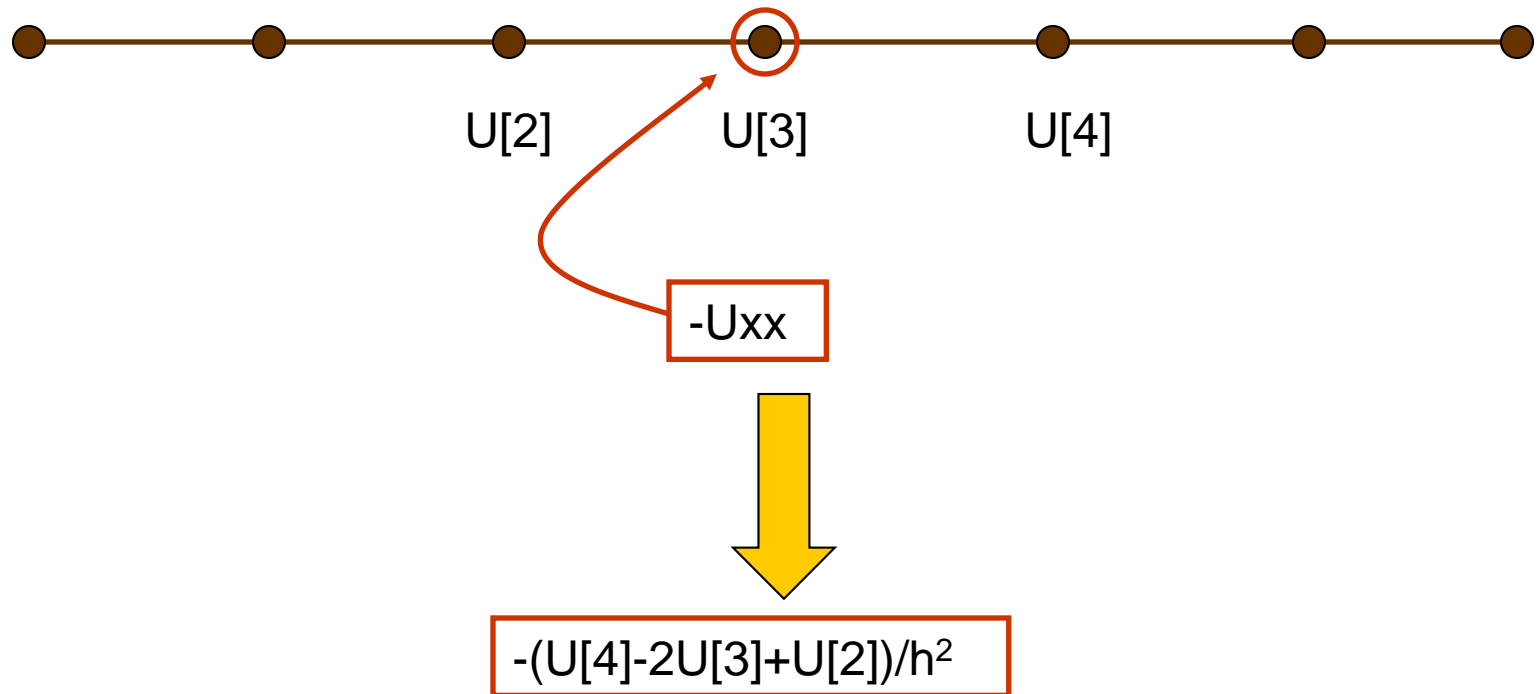
$$U(x-h) = U(x) - U'(x)h + \frac{1}{2}U''(x)h^2 - \frac{1}{6}U'''(x)h^3 + O(h^4)$$

□ Adding and solving for second derivative

$$U(x+h) + U(x-h) = 2U(x) + U''(x)h^2 + O(h^4)$$


$$U''(x) = \frac{1}{h^2} [U(x+h) - 2U(x) + U(x-h)] + O(h^2)$$

Discretizing: going from the continuum (PDE) to the discrete (algebraic)




Forward Difference method, replace derivatives by differences

$$-\frac{\partial^2 U(x)}{\partial x^2} = 1;$$




$$-\frac{U[i-1] - 2U[i] + U[i+1]}{h^2} = 1$$




Forward Difference method, replace derivatives by differences

$$-\frac{\partial^2 U(x)}{\partial x^2} = 1;$$




$$-\frac{U[i-1] - 2U[i] + U[i+1]}{h^2} = 1$$




Solve for U[i]

$$U[i] = \frac{1}{2} [h^2 + U[i-1] + U[i+1]]$$




Jacobi method, guess at a solution, and update via finite difference formula

$$U[i] = \frac{1}{2} [h^2 + U[i-1] + U[i+1]]$$



$$U_{new}[i] = \frac{1}{2} [h^2 + U_{old}[i-1] + U_{old}[i+1]]$$




Jacobi method, guess at a solution, and update via finite difference formula

Keep going ...

- $U[i] =$
- Replace U_{old} by U_{new}
 - Update again (called an iteration or step)
- If this process converges (U stops changing), we'll have ...

$$U_{new}[i] = \frac{1}{2} [h^2 + U_{old}[i-1] + U_{old}[i+1]]$$



Jacobi method, guess at a solution, and update via finite difference formula


Keep going ...

- Replace U_{old} by U_{new}
- Update again

If this process converges (U stops changing), we'll have a solution

$U[i] =$

$$U[i] = \frac{1}{2} [h^2 + U[i-1] + U[i+1]]$$



Serial Code Fragment: Jacobi

```
double          u_new[N+2], u_old[N+2];

u_old[0]=0.0, u_old[N+1]=0.0;
u_new[0]=0.0, u_new[N+1]=0.0;

h=1.0/(N+1), h2=h*h;

for (step=0; step<num_steps; step++) {
    for (i=1; i<=N; i++) {
        /* compute new temp from formula */
        u_new[i]=0.5*(h2+u_old[i+1]+u_old[i-1]);
    }
    for (i=1; i<=N; i++) {
        u_old[i]=u_new[i];
    }
}
```

(III) Heat Equation by Finite Differences + OpenMP



OpenMP is one standard for shared memory programming.

- ❑ Write most of code in C, C++, FORTRAN but add OpenMP compiler directives to control threads.
- ❑ Allows incremental parallelization
 - Profile serial code
 - Mark for parallelization those loops that take the most time
- ❑ Still have to think to make sure marked loops can be executed in parallel.

Jacobi Open MP

```
# include <omp.h>
double      u_new[N+2], u_old[N+2];
u_old[0]=0.0, u_old[N+1]=0.0;
u_new[0]=0.0, u_new[N+1]=0.0;
h=1.0/(N+1), h2=h*h;

for (step=0; step<num_steps; step++) {
#pragma omp parallel for
    for (i=1; i<=N; i++) {
        /* compute new temp from formula */
        u_new[i]=0.5*(h2+u_old[i+1]+u_old[i-1]);
    }
#pragma omp parallel for
    for (i=1; i<=N; i++) {
        u_old[i]=u_new[i];
    }
}
```


Jacobi Open MP

```
# include <omp.h>
```

```
double      u_new[N+2], u_old[N+2];
```

```
u_old[0]=0.0, u_old[N+1]=0.0;
```

```
u_new[0]=0.0, u_new[N+1]=0.0;
```

```
h=1.0/(N+1), h2=h*h;
```

Instructs the compiler to execute the immediately following for loop in parallel.

```
for (step=0; step<num_steps; step++) {
```

```
#pragma omp parallel for
```

```
for (i=1; i<=N; i++) {
```

```
/* compute new temp from formula */
```

```
u_new[i]=0.5*(h2+u_old[i+1]+u_old[i-1]);
```

```
}
```

```
#pragma omp parallel for
```

```
for (i=1; i<=N; i++) {
```

```
u_old[i]=u_new[i];
```

```
}
```

```
}
```

OpenMP: parallel for

- ▣ To allow compiler parallelize the loop, control clause must have canonical shape.

```
for(i = start; i      < end;      i++)
                >      ++i
                <=     i--
                >=     --i
                        i = i - inc
                        i -= inc
                        i = i + inc
                        i += inc
```

OpenMP: parallel for

- ▣ To allow compiler parallelize the loop, loop body can't contain statements that allow loop to exit prematurely.
 - No break
 - No return
 - No exit
 - No goto statements to labels outside the loop

OpenMP: parallel for

- ❑ Only mark as parallel a loop if there are no dependencies.
- ❑ The results should not depend on the order in which the loop is executed

OpenMP: parallel for

- ❑ Only mark as parallel a loop if there are no dependencies.
- ❑ The results should not depend on the order in which the loop is executed

```
for (i = 0; i < n; i++)  
{  
    c[i] = A[i] + B[i];  
}
```

OK to parallelize
Loop index k's computation is
independent of all others

OpenMP: parallel for

- ❑ Only mark as parallel a loop if there are no dependencies.
- ❑ The results should not depend on the order in which the loop is executed

```
for (i = 1; i < n; i++)  
{  
    c[i] = A[i] + c[i-1];  
}
```

Not OK to parallelize
Loop index k's computation is
dependent on result from k-1

Internally the loop iterations are divided among threads

```
#pragma omp parallel for
for (i = 0; i < n; i++)
{
    C[i] = A[i] + B[i];
}
```

OpenMP compiler will generate code like below:

```
int this_thread = omp_get_thread_num()
int num_threads = omp_get_num_threads();
int my_start = (this_thread ) * n / num_threads;
int my_end   = (this_thread+1) * n / num_threads;
for (i = my_start; i < my_end; i++)

{
    C[i] = A[i] + B[i];
}
```

OpenMP: shared and private variables

- A ***shared variable*** has the same address in every thread (there's only one version)
 - All threads can access shared variables
- A ***private variable*** has a different address in each thread (there's a version for each thread)
 - A thread cannot access a private variable of another thread
- Default for the **parallel for** pragma
 - All variables are shared except for the loop index which is private.

Jacobi Open MP

```
# include <omp.h>
```

```
double      u_new[N+2], u_old[N+2];
```

```
u_old[0]=0.0, u_old[N+1]=0.0;
```

```
u_new[0]=0.0, u_new[N+1]=0.0;
h=1.0/(N+1), h2=h*h;
```

Each thread will get its own, private "i" variable



```
for (step=0; step<num_steps; step++) {
```

```
#pragma omp parallel for
```

```
    for (i=1; i<=N; i++) {
```

```
        /* compute new temp from formula */
```

```
        u_new[i]=0.5*(h2+u_old[i+1]+u_old[i-1]);
```

```
    }
```

```
#pragma omp parallel for
```

```
    for (i=1; i<=N; i++) {
```

```
        u_old[i]=u_new[i];
```

```
    }
```

```
}
```

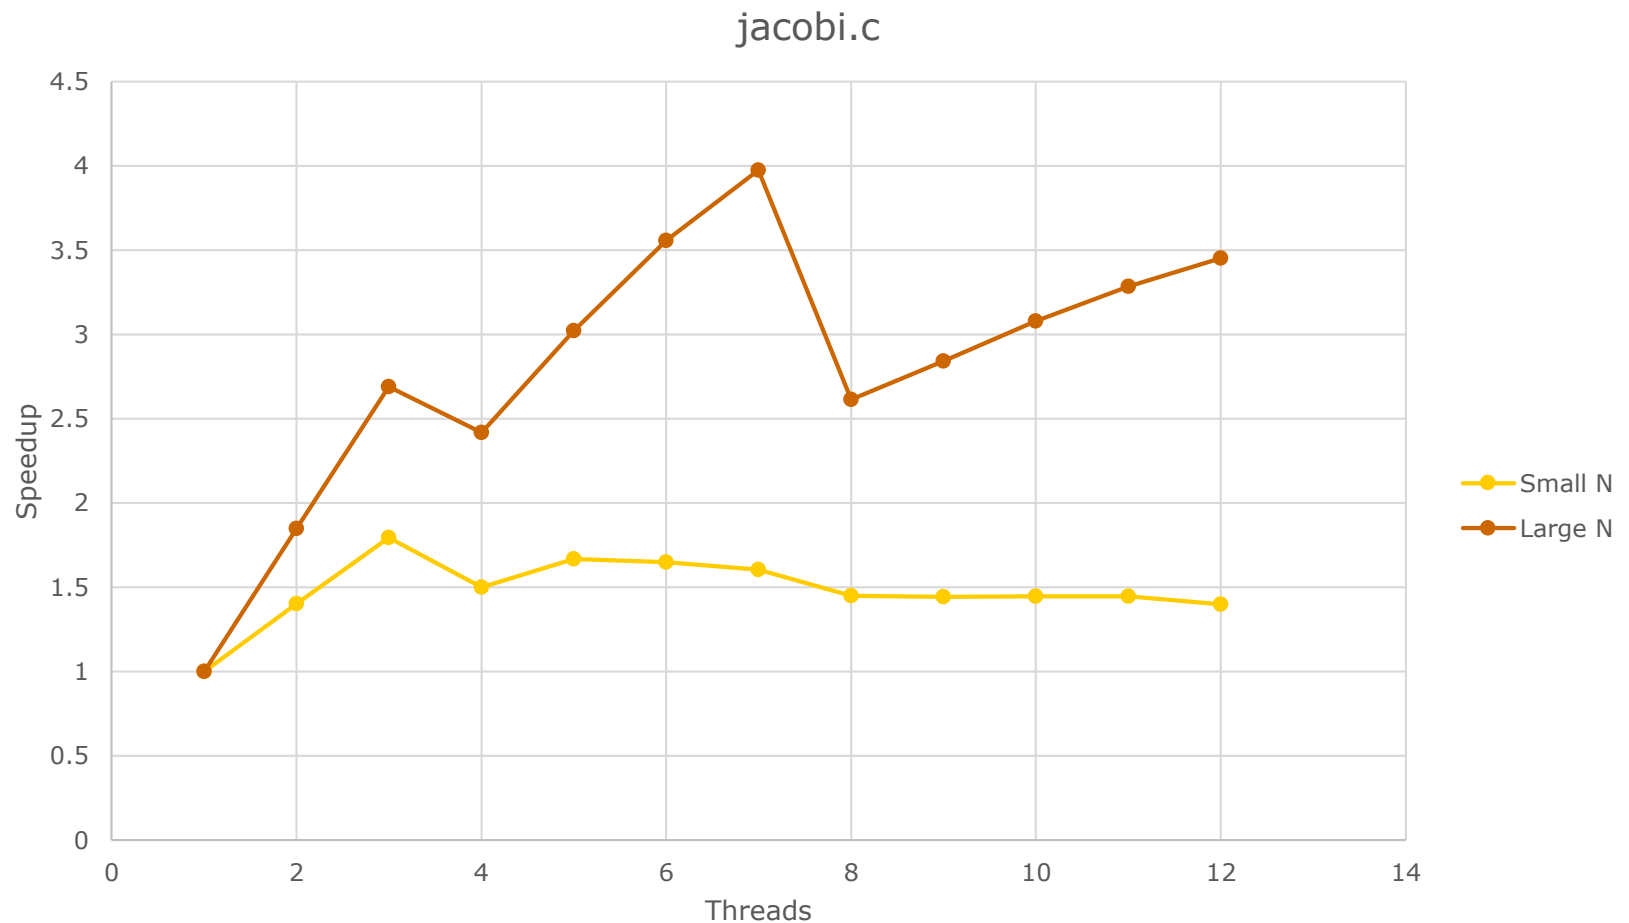
OpenMP: how many threads to use?

- ❑ `void omp_set_num_threads (int t)`
 - Sets the number of threads to be used in parallel sections.
 - Can also be controlled by the environment variable `OMP_NUM_THREADS`

Runtime (sec) for Jacobi OpenMP code on my quad core desktop

OMP num threads	Small N=16000	Large N=160000
1	16.16	182.88
2	11.54	98.92
3	9.01	67.99
4	10.77	75.63
5	9.69	60.49
6	9.80	51.42
7	10.07	46.00
8	11.16	69.97
9	11.2	64.36
10	11.18	59.40
11	11.17	55.67
12	11.55	52.99

Runtime (sec) for Jacobi OpenMP code on my quad core desktop



(IV) OpenMP



Codes on CANVAS

Parallel pragma alone

- ❑ The parallel pragma starts a parallel region.
- ❑ This starts (forks) a team of threads all of which execute the region. Each thread assigned an id: 0,1, ..., num_threads-1
- ❑ Implicit barrier at end of parallel region, threads wait until all finish
- ❑ After the region, threads join back to one.

```
#pragma omp parallel
{
    printf("Hello!\n");
}
```

Output:
One print from each thread