

Upcoming Schedule

- One more programming assignment, details likely Thursday after the break. Due date will be April 11, the last class meeting before final projects.
- Final Project presentations last 2 weeks of classes and maybe Finals week (Tue April 30 6pm).
- Proposals for projects will be due March 21.

Final Project preview

Proposal Due March 21

- Statement of the problem you want to solve.
- Outline of approach or algorithm you plan to use, References?
- Questions or issues do you plan to investigate.
- It want, you can discuss your idea with me beforehand.
- Undergrad projects can be a teamed project, no more than 3 students per team. Project scope should match team size.

Possibilities

- Numerical PDEs
- Numerical Linear Algebra
 - Sparse Matrix times Vector
- Monte Carlo Methods
 - Neutron Transport
- Sorting
- Others: searching, dense linear algebra, maximal independent sets, graph coloring ...
- Any topic from texts that we haven't covered
- Other styles of parallelism on other machines: “Proof of concept” recommended before proposal.

Numerical PDEs

Parallel multigrid solvers for partial differential equations in coastal ocean simulation

Jim E. Jones: Math

Steven Jachec: Marine and Environmental Systems

Anjali Ram PhD Math 2011


Osita Onyejekwe MS Math 2012

Work partially supported by NSF Grant No. CNS 09-23050
ONR Early Student Support Grant (# N000141110170)

Tidal effects: surface and internal

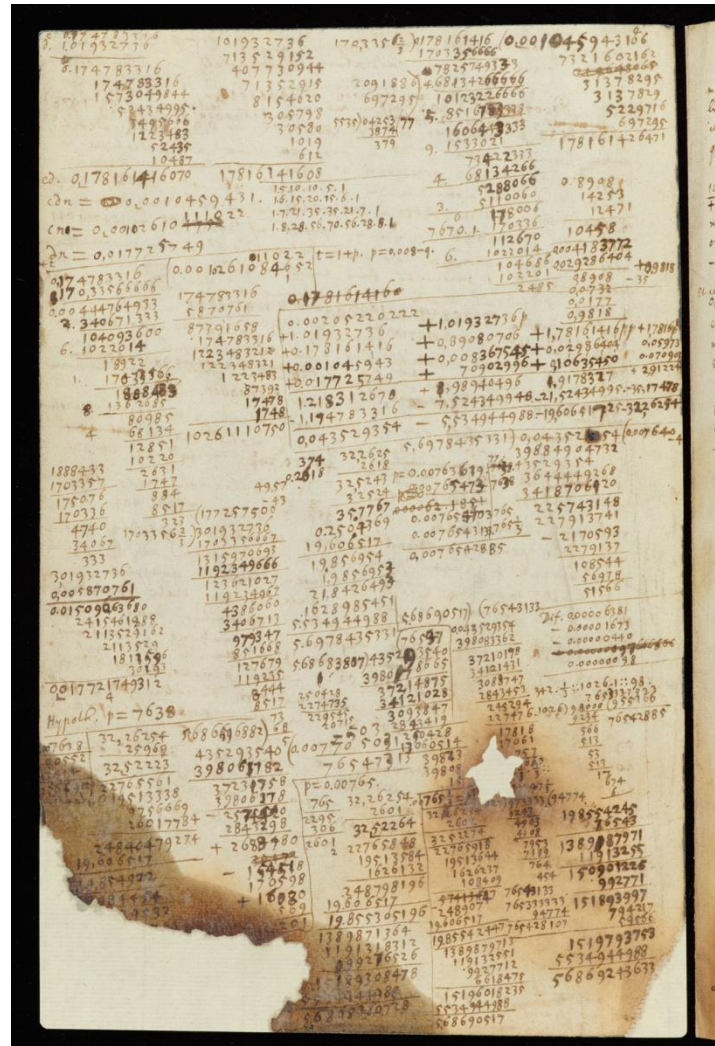
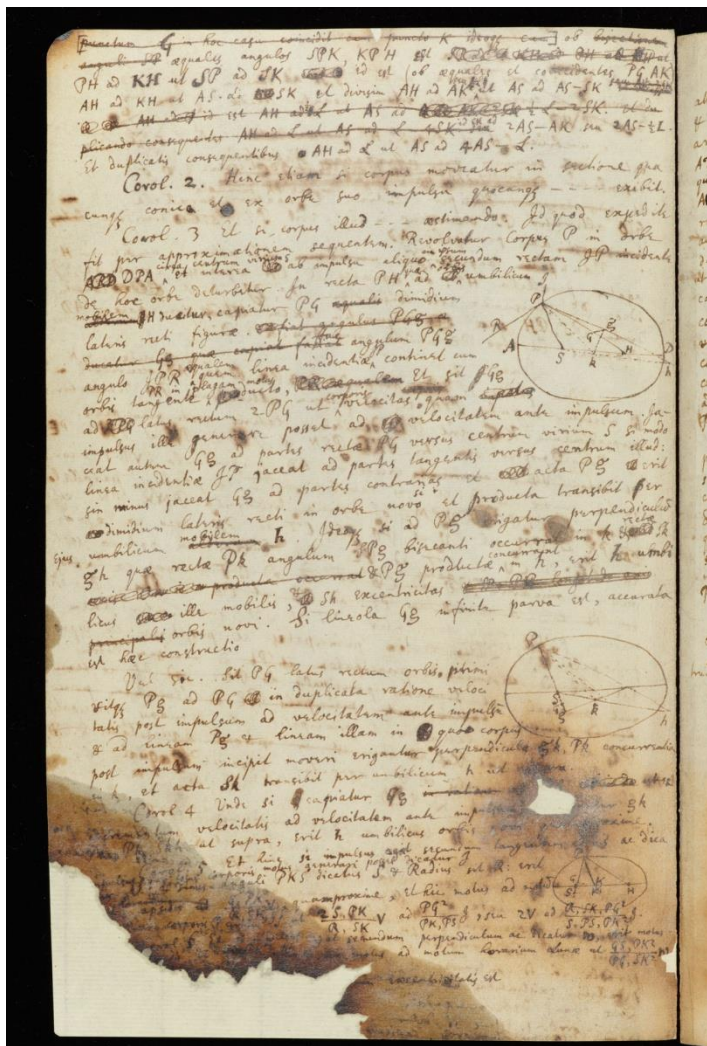


Bay of Fundy, New Brunswick, Canada




Pineda Woods Hole, Annual Report 1998

Early numerical model of tides



$$F = ma$$

$$F = G \frac{mM}{r^2}$$




Stanford Unstructured Non-hydrostatic Terrain-following Adaptive Navier- Stokes Simulator

The three-dimensional Navier-Stokes equations with the Boussinesq approximation in a rotating frame, after filtering with either Reynolds-averaging or via a large-eddy simulation and employing an eddy viscosity model, are given by

$$\frac{\partial u}{\partial t} + \nabla \cdot (\vec{u}u) - fv + bw = -\frac{1}{\rho_0} \frac{\partial p}{\partial x} + \nabla_H \cdot (\nu_H \nabla_H u) + \frac{\partial}{\partial z} (\nu_V \frac{\partial u}{\partial z})$$

$$\frac{\partial v}{\partial t} + \nabla \cdot (\vec{u}v) + fu = -\frac{1}{\rho_0} \frac{\partial p}{\partial y} + \nabla_H \cdot (\nu_H \nabla_H v) + \frac{\partial}{\partial z} (\nu_V \frac{\partial v}{\partial z})$$

$$\frac{\partial w}{\partial t} + \nabla \cdot (\vec{u}w) - bu = -\frac{1}{\rho_0} \frac{\partial p}{\partial z} + \nabla_H \cdot (\nu_H \nabla_H w) + \frac{\partial}{\partial z} (\nu_V \frac{\partial w}{\partial z}) - \frac{g}{\rho_0} (\rho_0 + \rho)$$



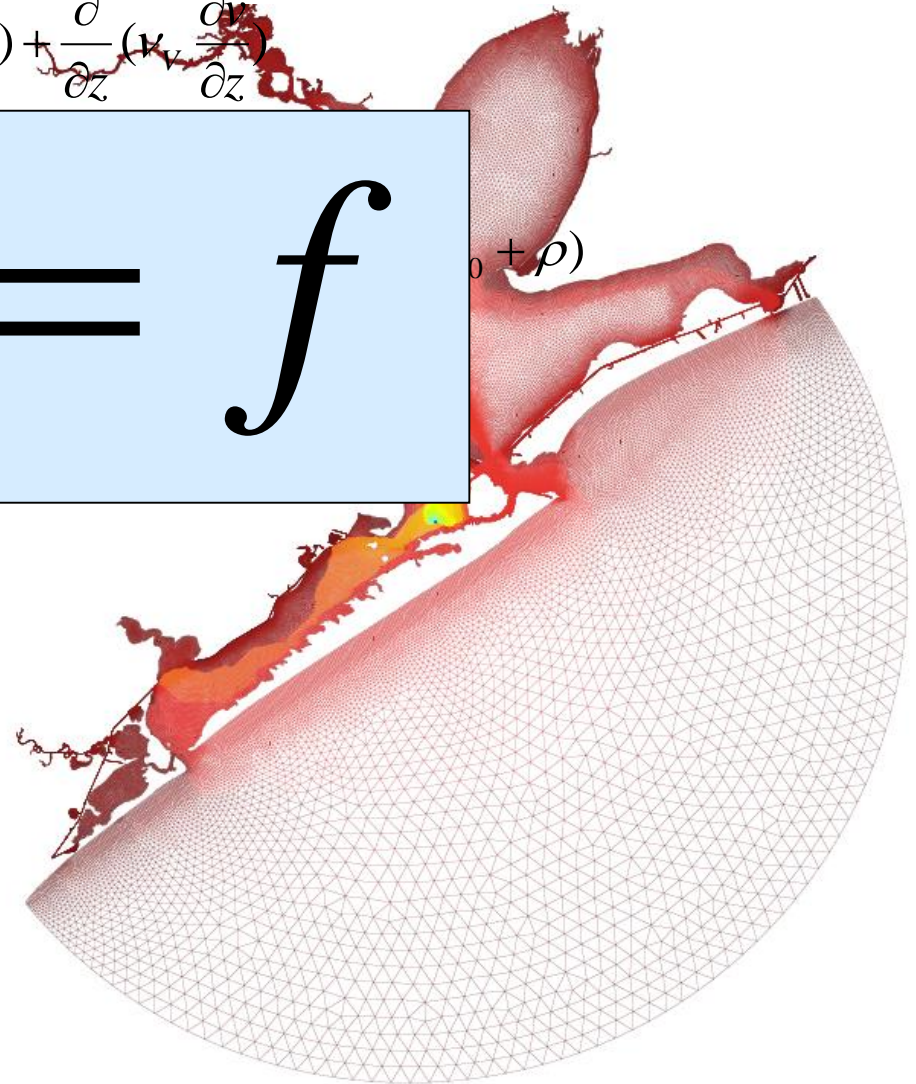

<http://www.stanford.edu/~fringer/research.html>

$$\frac{\partial u}{\partial t} + \nabla \cdot (\vec{u}u) - fv + bw = -\frac{1}{\rho_0} \frac{\partial p}{\partial x} + \nabla_H \cdot (v_H \nabla_H u) + \frac{\partial}{\partial z} (v_v \frac{\partial u}{\partial z})$$

$$\frac{\partial v}{\partial t} + \nabla \cdot (\vec{u}v) + fu = -\frac{1}{\rho_0} \frac{\partial p}{\partial y} + \nabla_H \cdot (v_H \nabla_H v) + \frac{\partial}{\partial z} (v_v \frac{\partial v}{\partial z})$$

$$\frac{\partial w}{\partial t}$$



$$Au = f$$



If you can't do a problem, there is a simpler problem you do not understand. Your first job is to find that simpler problem.


$$u_t - (u_{xx} + u_{yy}) = f$$

$$Au = f$$



Discretization approximates the differential problem by an algebraic one


$$u_{xx} + u_{yy}$$




$$Au = f$$

Continuous to Discrete in 2D

$$u_{xx}(x, y) \approx \frac{1}{h^2} (u(x-h, y) - 2u(x, y) + u(x+h, y))$$




$$u_{xx}(x, y) \approx \frac{1}{h^2} (u(x-h, y) - 2u(x, y) + u(x+h, y))$$



Approximated derivative at a point by an algebraic equation involving function values at nearby points

By Taylor's Theorem, the error in this approximation (the truncation error) is **$O(h^2)$**

$$u_{xx} + u_{yy} \quad \longrightarrow \quad \frac{1}{h^2} (u(x-h, y) + u(x, y-h) - 4u(x, y) + u(x+h, y) + u(x, y+h))$$



Error in approximation is determined by the mesh size ***h***. Difference between differential solution and algebraic solution goes to zero as ***h*** does.

Solve the heat equation in 2d

- Develop Parallel Code.
- Run on Blueshark.
- Investigate speed-ups and scalability.
- Explore different partition strategies.

Related PDE possible projects

- Implicit methods for time dependent problems.
- Use parallel libraries (hypre, PetSc, ...)
- Run on Blueshark.
- Investigate speed-ups and scalability.

Sparse MPI Matvec

Sparse Matrix times Dense Vector

- $y = Ax$
- Store only non-zero's of A

$$\begin{bmatrix} 6 & 0 & 0 & 4 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2 & 8 & 0 & 0 & 0 & 0 & 0 & 2 & 0 & 0 \\ 0 & 0 & 4 & 0 & 0 & 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 2 & 2 & 0 & 9 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 10 & 0 & 15 & 0 & 0 & 1 & 0 \\ 0 & 2 & 0 & 0 & 0 & 0 & 5 & 0 & 0 & 1 \\ 2 & 0 & 0 & 0 & 0 & 1 & 0 & 8 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 10 & 0 \\ 0 & 12 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 37 \end{bmatrix} \begin{bmatrix} 2 \\ 8 \\ 5 \\ 6 \\ 12 \\ 44 \\ 1 \\ 9 \\ 0 \\ 14 \end{bmatrix}$$


Sparse Matrix times Dense Vector

- Develop Parallel MPI Code.
- Challenge is minimize the communication and allow general sparsity structure.
- Run on Blueshark.
- Investigate speed-ups and scalability.

Monte Carlo

Monte Carlo Methods


- Ratio of areas



$$\frac{\pi D^2 / 4}{D^2} = \frac{\pi}{4}$$


Monte Carlo Methods

- Ratio of areas



- Ratio of “hits”

Monte Carlo Methods



- Monte Carlo approximates
- Parallelization by running trials on each processor

Monte Carlo Methods

- Issues

- accuracy vs. number of trials
- parallel random number generators

- Applications

- Neutron Transport

Monte Carlo Methods


□ Issues

- accuracy vs. number of trials
- parallel random number generators

□ Applications

- Neutron Transport

$$C = C_s + C_c$$



While not done

Travel Random Distance $L = -\ln(u)/C$

If outside slab, done.

Interact with atom (random outcome)

If $(u < C_c/C)$ absorbed, done.

Else scattered in random direction ($d = 2\pi \cdot u$)

Monte Carlo Methods

- Pick an application (Neutron transport?)
- Develop Parallel Code.
 - Will require mostly looking at parallel random number generation.
- Run on Blueshark.
- Investigate speed-ups and scalability.
- Investigate accuracy vs. number of trials

Sorting

Parallel Sorting

- Given a set of n integers on each processor

| | | | |
|---|----|---|----|
| 6 | 20 | 5 | 73 |
|---|----|---|----|

| | | | |
|---|-----|---|----|
| 1 | 101 | 3 | 66 |
|---|-----|---|----|

| | | | |
|---|----|-----|----|
| 1 | 32 | 261 | 55 |
|---|----|-----|----|

- Produce a sorted list: sorted within a processor and largest integer on p_0 less than or equal to smallest on p_1 when $\text{rank}(p_0) < \text{rank}(p_1)$.

| | | | |
|---|---|---|---|
| 1 | 1 | 3 | 5 |
|---|---|---|---|

| | | | |
|---|----|----|----|
| 6 | 20 | 32 | 55 |
|---|----|----|----|

| | | | |
|----|----|-----|-----|
| 66 | 73 | 101 | 261 |
|----|----|-----|-----|

Parallel Sorting

- Quinn Ch 14 discusses Parallel Sorting.
- I can provide other references as well.
- Project might be
 - Implement one or more of the other parallel sorting algorithms
 - Compare to Sorting performance.
 - Compare to models.
- FYI: this is a popular project topic but is difficult to do well and see significant speedups.


Worker Manager

- Certain applications will be well suited to this style:
 - Game trees
 - Number theory problems
 - Any problem where tasks vary in length in unpredictable ways
- Run on Blueshark.
- Investigate speed-ups and scalability.

Graph Coloring


Graph Theory Definitions

- A **graph** $G=(V,E)$ is a collection of **vertices** $V=\{v_1, v_2, \dots, v_n\}$ connected by **edges** $E=\{e_1, e_2, \dots, e_m\}$.




Independent Set Definition

- A set of vertices is a **independent set** if none of the vertices are connected by an edge




Maximal Independent Set Definition

- A set S of vertices is a **maximal independent set** (MIS) if it is an independent set and adding any other vertex to the set produces a set which is not independent.




Independent Set Example

- $S_1 = \{1, 5\}$ is not a MIS because we could add, say, vertex 7 to it and still have an independent set.




Maximal Independent Set Example

- $S_2 = \{1, 3, 10\}$ is a MIS




Maximal Independent Set is not unique

- $S_3 = \{4, 6, 8, 10\}$ is a MIS




Maximal Independent Set Algorithm?

- Given a graph $G=(V,E)$, find a maximal independent set S .



Maximal Independent Set Algorithm?

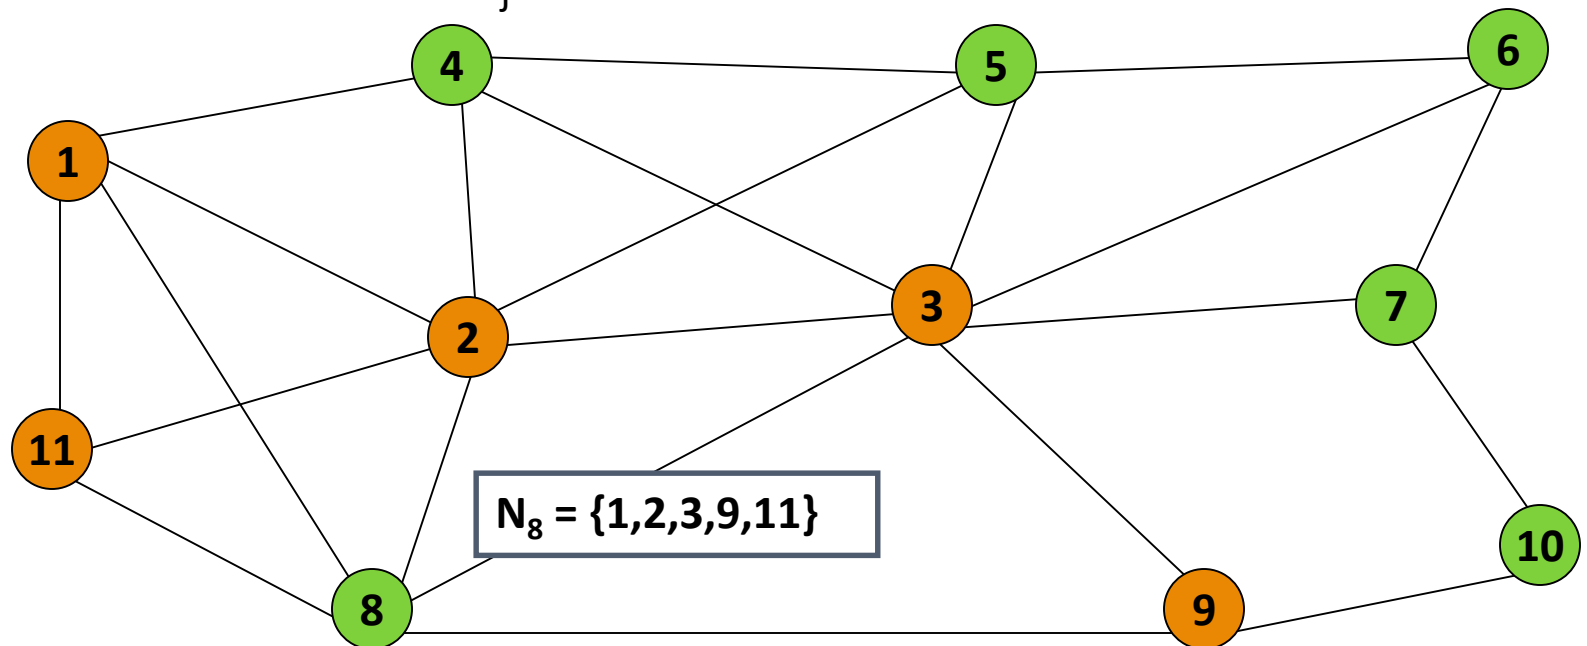
- For now, assume one process owns the entire graph.
Ideas???????



Maximal Independent Set Algorithm?

Hint # 1

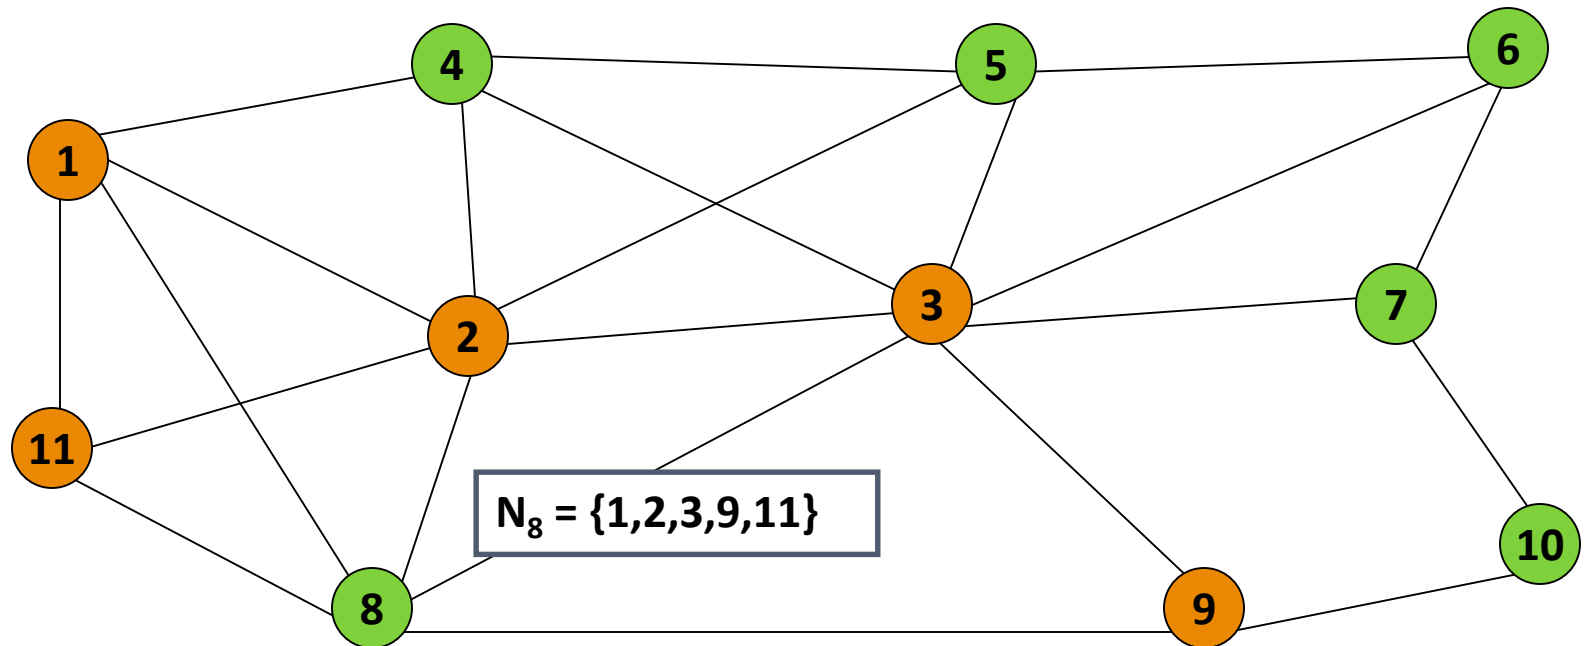
- The **neighborhood** N_j of a vertex v_j consists of all vertices connected to v_j by an edge.
- The number of vertices in the neighborhood is called the **degree** of vertex v_j



Maximal Independent Set Algorithm?


Hint # 2

- If vertex v_j is in set S and S is a maximal independent set then all vertices in neighborhood N_j are not in S .



Maximal Independent Set Algorithm

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Pick v_j in U and add to S
 2. Remove v_j and its neighborhood N_j from U




Maximal Independent Set Algorithm

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Pick v_j in U and add to S
 2. Remove v_j and its neighborhood N_j from U

$U = \{1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11\}$

$S = \{ \}$




Maximal Independent Set Algorithm

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Pick v_j in U and add to S
 2. Remove v_j and its neighborhood N_j from U

$U = \{1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11\}$

$S = \{1\}$




Maximal Independent Set Algorithm

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Pick v_j in U and add to S
 2. Remove v_j and its neighborhood N_j from U

$U = \{3, 5, 6, 7, 9, 10\}$

$S = \{1\}$

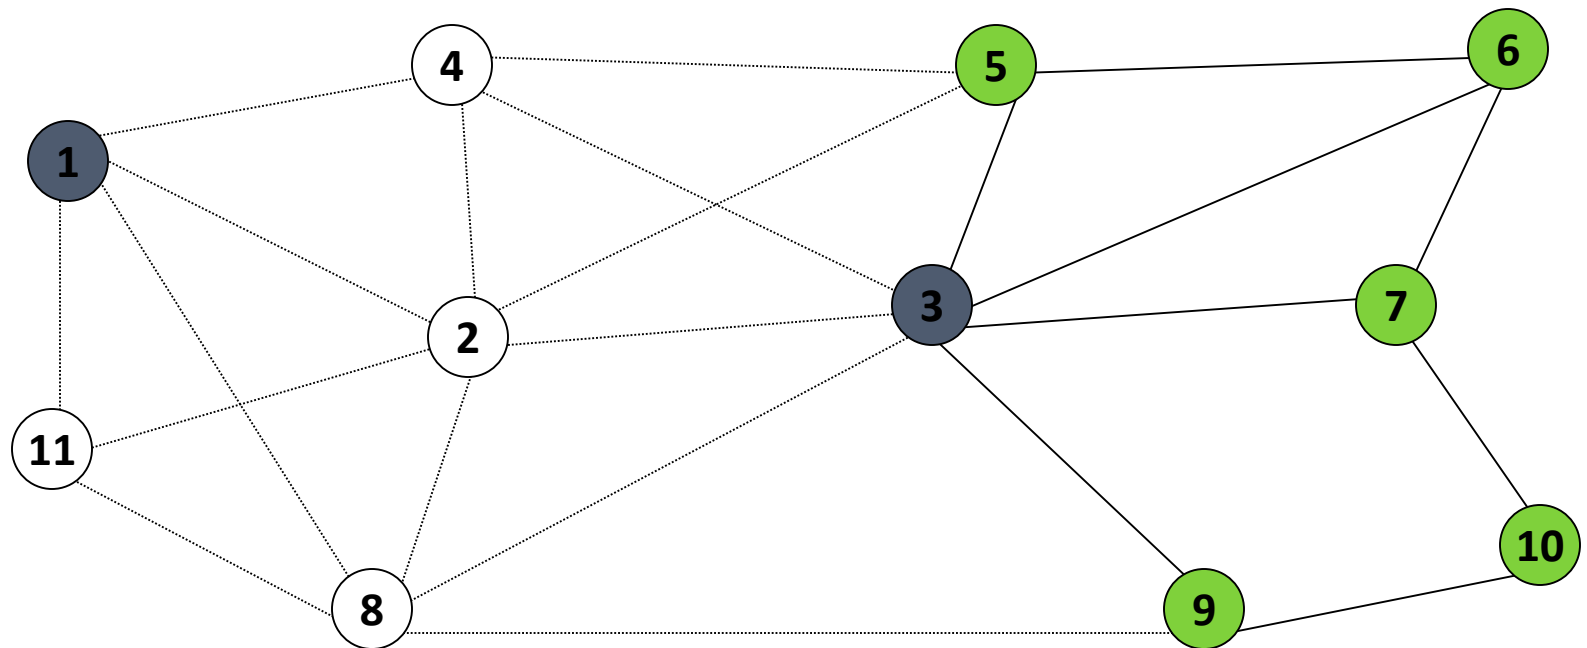


Maximal Independent Set Algorithm

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Pick v_j in U and add to S
 2. Remove v_j and its neighborhood N_j from U

$U = \{3, 5, 6, 7, 9, 10\}$

$S = \{1, 3\}$




Maximal Independent Set Algorithm

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Pick v_j in U and add to S
 2. Remove v_j and its neighborhood N_j from U

$U = \{10\}$

$S = \{1, 3\}$




Maximal Independent Set Algorithm

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Pick v_j in U and add to S
 2. Remove v_j and its neighborhood N_j from U

$U = \{10\}$

$S = \{1, 3, 10\}$




Maximal Independent Set Algorithm

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Pick v_j in U and add to S
 2. Remove v_j and its neighborhood N_j from U

$U = \{ \}$


$S = \{1, 3, 10\}$

Done as U now empty




Maximal Independent Set Parallel Algorithm?

- Given a graph $G=(V,E)$, find a maximal independent set S .
- Assume the number of processes is equal to n , the number of vertices.



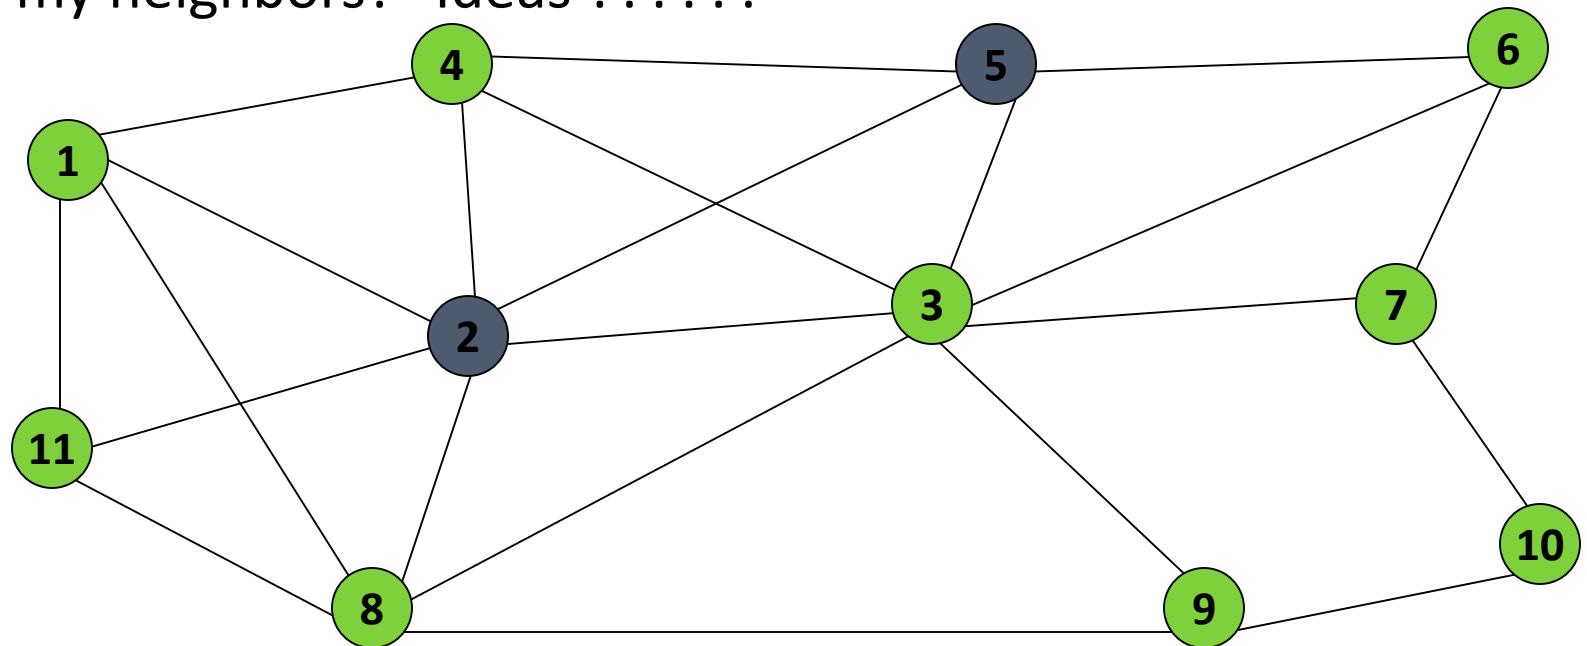
Maximal Independent Set Parallel Algorithm?

- Choosing if my vertex is in S or not is a candidate for the primitive task.
- Hurdle: how do I make sure my choice doesn't conflict with my neighbors?



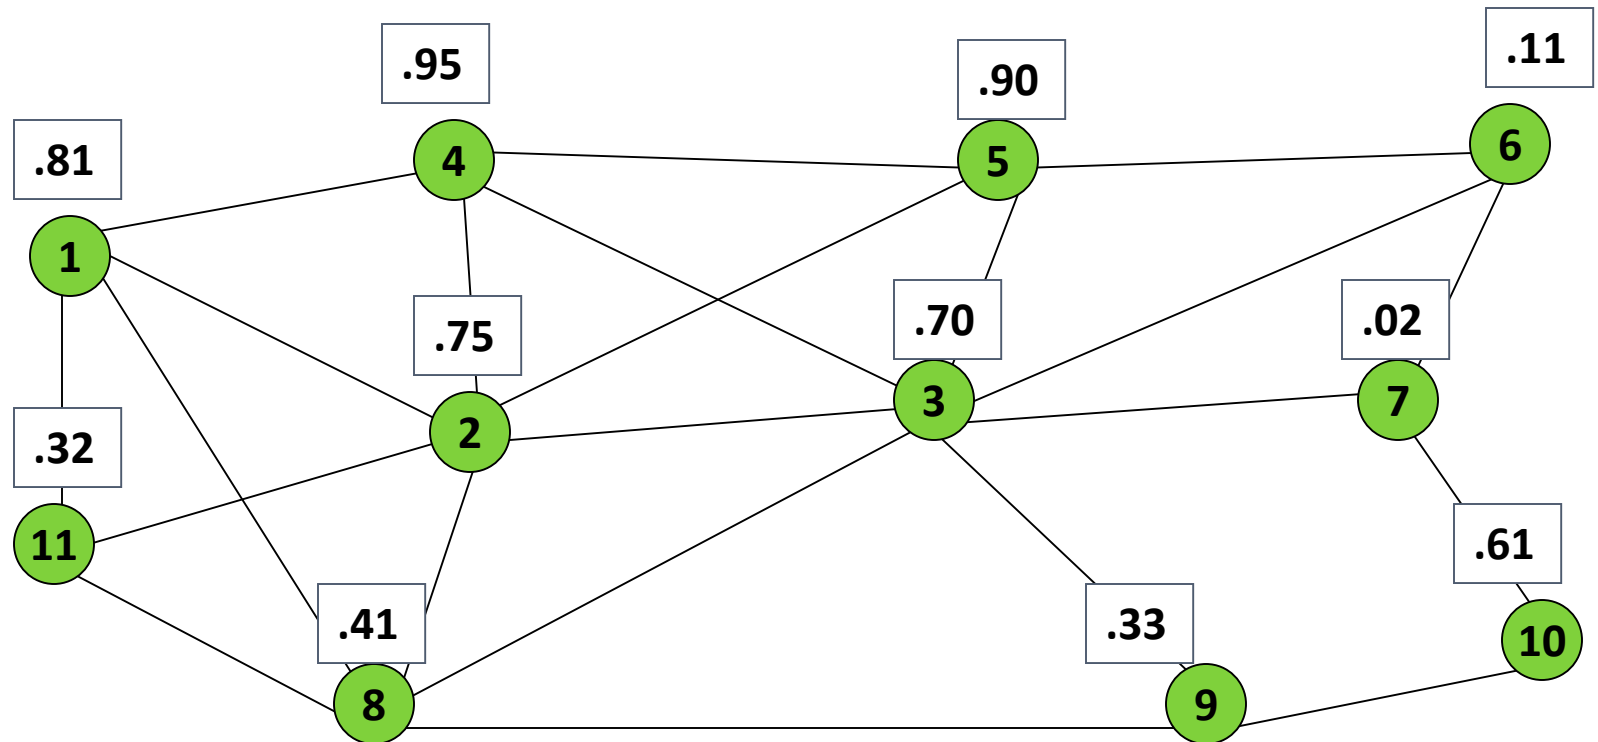
Maximal Independent Set Parallel Algorithm?

- Choosing if my vertex is in S or not is a candidate for the primitive task.
- Hurdle: how do I make sure my choice doesn't conflict with my neighbors? Ideas ??????



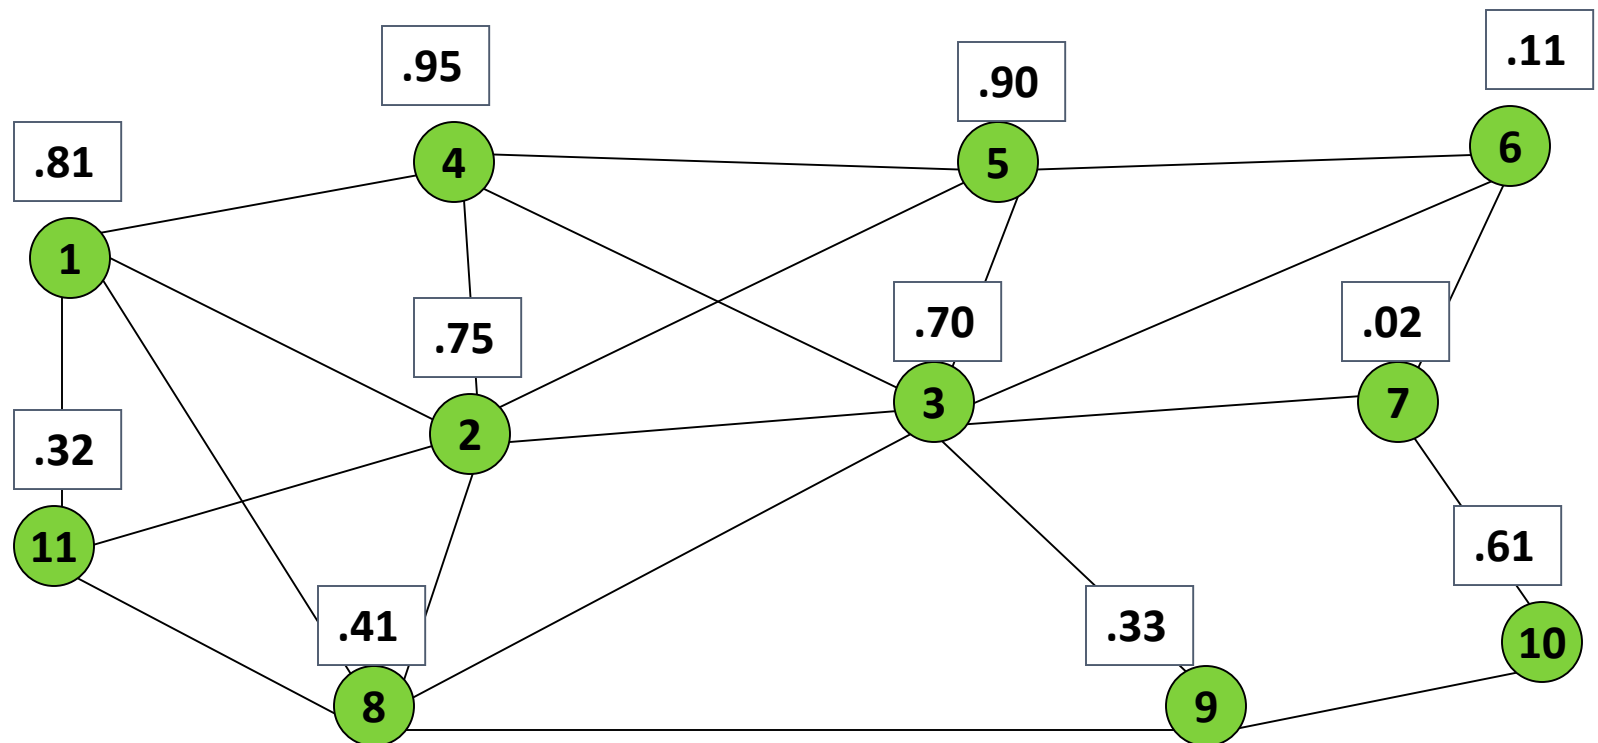
Maximal Independent Set Parallel Algorithm? Hint #1

- Assign a random number r_i in $(0,1)$ to my vertex. Communicate this value to neighbors.



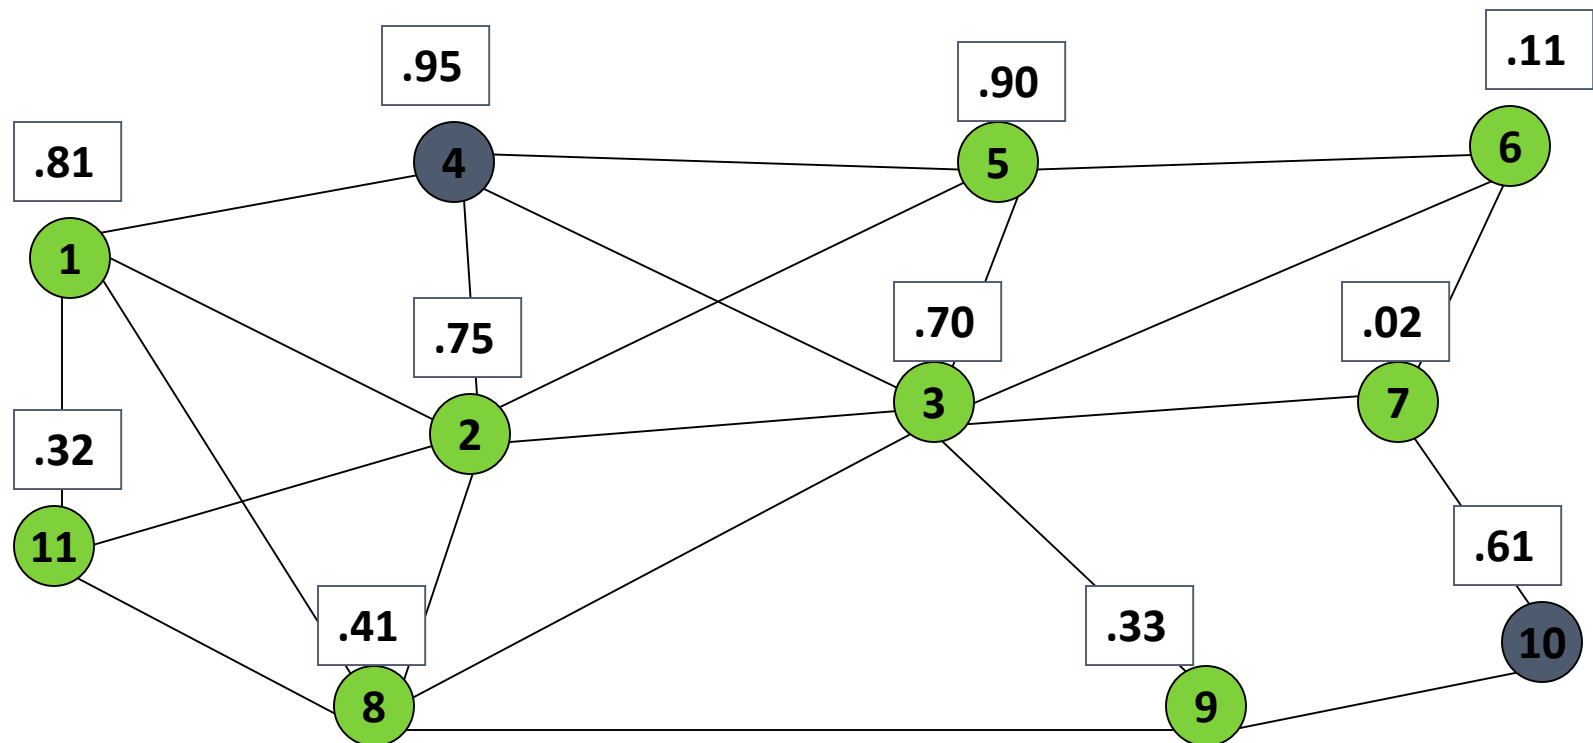
Maximal Independent Set Parallel Algorithm? Hint #2

- If my r is bigger than all my neighbor's r I'm in S .



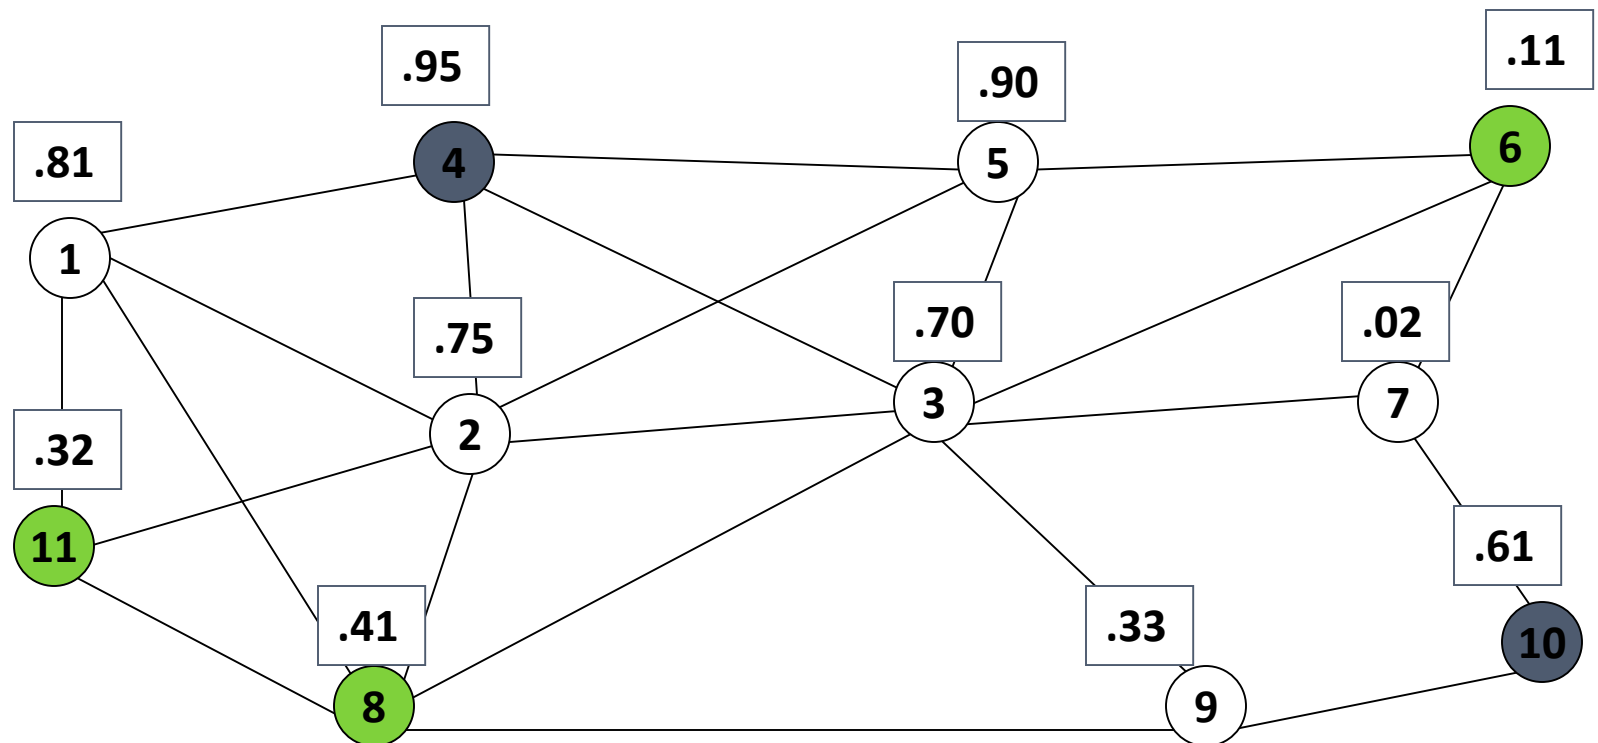
Maximal Independent Set Parallel Algorithm? Hint #2

- If my r is bigger than all my neighbor's r , I'm in S .



Maximal Independent Set Parallel Algorithm? Hint #3

- Remove vertices and edges for everyone in S and their neighbors (requires communication).



Maximal Independent Set Parallel Algorithm? Hint #3

- Start over with remaining graph.




Maximal Independent Set Parallel Algorithm? Hint #3

- Start over with remaining graph.



Maximal Independent Set Algorithm

- No more undecided vertices.



Maximal Independent Set Algorithm:


Luby: SIAM J. Comput., 15(4):1036--1053

1. $S = \text{empty}$, $U = V$
2. While U not empty
 1. Assign random r_i in $(0,1)$ to each vertex
 2. If r_j is greater than neighbors add v_j to S^*
 3. Add S^* to S
 4. Remove S^* and its neighbors from U

- Which MIS we get depends on the random numbers.
- Number of iterations of while loop depends on random numbers.


Maximal Independent Set Parallel Algorithm?

- Given a graph $G=(V,E)$, find a maximal independent set S .
- Assume the number of processes is much smaller than n , the number of vertices.



Maximal Independent Set Parallel Algorithm?

- Given a graph $G=(V,E)$, find a maximal independent set S .
- Assume the number of processes is much smaller than n , the number of vertices.
- The graph is distributed. Ideas?



Graph Coloring

- Develop Parallel Code.
- Run on Blueshark.
- Investigate speed-ups and scalability.

Possibilities

- Numerical PDEs
- Numerical Linear Algebra
 - Sparse Matrix times Vector
- Monte Carlo Methods
 - Neutron Transport
- Sorting
- Others: searching, dense linear algebra, maximal independent sets, graph coloring ...
- Any topic from texts that we haven't covered
 - Pacheco: n-Body Solvers and Tree Search
 - Quinn: lots – sorting, searching, FFT, Matrix Algebra
- Other styles of parallelism on other machines: like GPU. “Proof of concept” recommended before proposal.
- Same originality “Rules” as assignments: the code your team turns in must be written by your team.